Job Title

Senior Data Scientist

Company Overview

Orion Commerce Analytics Pvt. Ltd. is a product-analytics company based in Pune. It provides pricing, demand forecasting and customer analytics platforms to e-commerce and consumer brands in India and Southeast Asia.

Role Summary

The Senior Data Scientist owns analytical solutions end to end: framing the business problem, designing experiments, building and validating models, and working with engineers to put them into monitored production use. The role also involves mentoring two to three junior data scientists and reporting business impact to senior stakeholders.

Key Responsibilities

- Frame ambiguous business problems as measurable data science problems with agreed success metrics.
- Design and analyse experiments, including A/B tests, and apply statistical modelling such as regression, time-series or survival methods where suitable.
- Develop supervised and unsupervised models (classification, regression, clustering) with documented feature engineering.
- Validate models rigorously using suitable cross-validation, leakage checks, calibration and error analysis.
- Work with machine learning and data engineers to deploy models as batch or API services and define monitoring for data drift and model performance.
- Quantify business impact of models and communicate results to senior stakeholders.
- Mentor junior data scientists through code review, analysis-plan review and day-to-day guidance.

Required Technical Skills

- Advanced Python (Pandas, NumPy, Scikit-learn, Statsmodels) and advanced SQL (CTEs, window functions, query optimisation).
- Statistical modelling and experimentation: hypothesis testing, A/B test design, power analysis and regression.
- Supervised and unsupervised machine learning: gradient boosting (XGBoost or LightGBM), random forests, clustering and dimensionality reduction.
- Feature engineering and model validation: cross-validation strategies, leakage prevention, metric selection and calibration.
- Working knowledge of model deployment and monitoring: packaging models, batch or API serving with engineering teams, drift monitoring and retraining triggers.
- Data visualisation and storytelling for business stakeholders.

Required Experience

- 5 to 8 years in data science or applied analytics.
- At least 3 years building predictive or statistical models that were used in business decisions.
- At least one model or analytical solution taken into production or regular operational use, with the candidate's own contribution clearly described.
- At least one example of measured business impact with a quantified outcome.
- At least 1 year of mentoring or reviewing the work of others, formal or informal.

Education

Bachelor's degree in Computer Science, Statistics, Mathematics, Engineering or a related field. A Master's degree (M.Tech, M.Sc., MCA) is preferred but not mandatory. Equivalent practical experience may substitute for formal qualifications.

Certifications

Mandatory certifications: None.

Preferred certifications (not mandatory): AWS Certified Machine Learning - Specialty, Microsoft Azure Data Scientist Associate (DP-100), Google Cloud Professional Machine Learning Engineer, or a recognised statistics or experimentation course certificate.

Preferred Qualifications

- Time-series forecasting.
- Deep learning or NLP.
- Spark or PySpark.
- MLflow, Airflow or Docker.
- Causal inference or uplift modelling.
- Retail or e-commerce domain experience.

Evaluation Priorities

- Models in production or regular operational use, with the candidate's own role and a quantified impact.
- Depth in experimentation and validation methods, described beyond a list of model names.
- Evidence of working with engineers on deployment and monitoring.
- Evidence of mentoring or technical leadership.
- Preferred technologies are secondary and should not outweigh the required evidence. Certifications are preferred only.

Employment Details

- Location: Pune, Maharashtra
- Employment type: Full-time, permanent
- Work arrangement: Hybrid, two days per week in office