Job Title

AI / Generative AI Engineer

Company Overview

Lumen Cloud Systems Pvt. Ltd. is an enterprise software company based in Hyderabad. It builds knowledge-management and workflow products for banks, insurers and manufacturers in India, and is adding generative AI features to its platform.

Role Summary

The AI / Generative AI Engineer designs, builds and deploys applications powered by large language models, such as document question answering, summarisation and assistants. The work centres on retrieval-augmented generation (RAG), prompt design, evaluation of answer quality and integration of AI services into production applications.

Key Responsibilities

- Build LLM-powered applications using hosted or open-source models through APIs.
- Build RAG pipelines covering document ingestion, chunking, embeddings, indexing in a vector database, retrieval, reranking and grounded answer generation.
- Write prompts and enforce structured outputs (JSON schema validation) with citations to source passages.
- Create test sets and measure retrieval quality (for example recall@k and MRR) and answer quality, and apply groundedness checks to reduce hallucination.
- Integrate AI services into applications through REST APIs, handling latency, caching, token cost and error handling.
- Apply responsible-AI practices: PII handling, access control on retrieved content, awareness of prompt injection and audit logging.
- Deploy services with Docker on a cloud platform and monitor them.

Required Technical Skills

- Python and REST API development (FastAPI or Flask).
- Hands-on LLM application development with hosted or open models, including prompt engineering and structured (JSON) outputs.
- RAG design and implementation: chunking, embeddings, vector search and retrieval, shown in at least one project.
- Embeddings and at least one vector database or index (for example FAISS, Chroma, Qdrant or pgvector).
- Evaluation of LLM and retrieval outputs: test sets, retrieval metrics, groundedness and hallucination checks.
- Git, Docker basics, JSON and basic SQL.
- Orchestration: experience with LangChain, LlamaIndex or a comparable framework, or with custom orchestration code. No specific framework is mandatory.

Required Experience

- 2 to 5 years in software engineering, machine learning or data roles.
- At least 1 year, or two substantial projects, building LLM or RAG applications.
- At least a working prototype with some form of evaluation. Production or pilot deployment is preferred but not mandatory.
- Experience limited to classical machine learning without LLM or RAG work does not meet the core requirement.

Education

B.Tech, B.E., MCA, M.Tech or M.Sc. in Computer Science, Information Technology, AI or a related field. Candidates with other qualifications may be considered when project evidence is strong.

Certifications

Mandatory certifications: None.

Preferred certifications (not mandatory): Microsoft Azure AI Engineer Associate (AI-102), Azure AI Fundamentals (AI-900), an AWS or Google Cloud certification, or a generative AI course certificate.

Preferred Qualifications

- Fine-tuning of language models (for example LoRA).
- Agent or tool-calling workflows.
- Hybrid search (BM25 plus dense retrieval) and rerankers.
- Kubernetes.
- Hosting of open-source LLMs.
- Guardrail libraries and LLM application security testing.
- Streamlit for internal demos.

Evaluation Priorities

- An end-to-end RAG or LLM application with named components: embeddings, vector store, retrieval and generation.
- Evaluation work, such as test sets, metrics and hallucination mitigation.
- API integration and deployment evidence.
- Specific framework names are secondary to evidence of the underlying concepts.
- Certifications are preferred only.

Employment Details

- Location: Hyderabad, Telangana
- Employment type: Full-time, permanent
- Work arrangement: Hybrid, three days per week in office