Job Title

Machine Learning Engineer

Company Overview

Nilgiri Fintech Technologies Pvt. Ltd. is a digital lending and payments company based in Chennai. Its machine learning platform team supports credit-risk, fraud and customer-engagement models used across its mobile and web products.

Role Summary

The Machine Learning Engineer turns prototype models into reliable production services. The role focuses on feature pipelines, model training and tuning, API serving, versioning, monitoring and reproducibility, working closely with data scientists who develop the analytical approach. It is an engineering role rather than a research role.

Key Responsibilities

- Build data preparation and feature engineering pipelines in Python and SQL that are reproducible and tested.
- Train, evaluate and tune models with Scikit-learn, XGBoost and a deep learning framework such as PyTorch or TensorFlow against defined metrics.
- Package and serve models through REST APIs (FastAPI or Flask) that meet latency and throughput targets.
- Build ML pipelines with experiment tracking, data and model versioning (for example MLflow or DVC) and scheduled retraining.
- Containerise services with Docker and deploy them on a cloud platform through CI/CD.
- Monitor deployed models for performance and data drift and set up alerts.
- Write tests and documentation, review code, and help data scientists convert prototype notebooks into production code.

Required Technical Skills

- Python (production-quality code) and SQL.
- Data preparation and feature engineering with Pandas; Spark is helpful.
- Model training, evaluation and tuning with Scikit-learn, plus working knowledge of at least one deep learning framework (PyTorch or TensorFlow).
- Hands-on model serving through REST APIs (FastAPI or Flask).
- ML pipeline practices: experiment tracking, data and model versioning and reproducibility (MLflow, DVC or equivalent).
- Docker and Git, plus hands-on use of at least one cloud platform (AWS, Azure or GCP) or a CI/CD tool.
- Understanding of model monitoring concepts, including drift and performance tracking.

Required Experience

- 3 to 6 years in machine learning engineering, software engineering with ML, or data science with production exposure.
- At least 2 years of deploying or maintaining ML models in production or production-like environments.
- At least one model deployed as an API service, with the candidate's own contribution clearly described.
- Purely analytical or research experience without deployment work does not meet the core requirement.

Education

B.Tech, B.E., M.Tech or MCA in Computer Science, Information Technology, Electronics or a related field.

Certifications

Mandatory certifications: None.

Preferred certifications (not mandatory): AWS Certified Machine Learning - Specialty, Microsoft Azure Data Scientist Associate (DP-100), Google Cloud Professional Machine Learning Engineer, or Certified Kubernetes Application Developer (CKAD).

Preferred Qualifications

- Kubernetes.
- Airflow or Kubeflow pipelines.
- Feature stores.
- Spark or PySpark.
- Model optimisation such as ONNX export.
- Prometheus and Grafana monitoring.
- Experience serving large language models.

Evaluation Priorities

- Production deployment of models, covering API serving, containers and cloud, with the candidate's own role.
- Evidence of reproducible pipelines, versioning and monitoring.
- Engineering practices such as testing and CI/CD.
- Competence in model development and evaluation.
- Certifications are preferred only.

Employment Details

- Location: Chennai, Tamil Nadu
- Employment type: Full-time, permanent
- Work arrangement: Hybrid, three days per week in office