# Farhan Qureshi

**Email:** john.aimldev@gmail.com
**Phone:** +91 00000 00111
**Location:** Hyderabad, Telangana, India

## Professional Summary

Software engineer turned GenAI engineer with five years of experience, the last two and a half building retrieval-augmented generation systems for enterprise knowledge and support use cases. Strong in backend design, embeddings, vector databases and LLM evaluation. Experience with traditional model training is limited.

## Technical Skills

**Programming:** Python, TypeScript (basic), SQL

**RAG and LLMs:** LangChain, LlamaIndex, Azure OpenAI, hybrid retrieval, re-ranking, chunking, prompt engineering, guardrails

**Vector Databases:** Qdrant, Pinecone, pgvector, FAISS

**Evaluation:** RAGAS, golden question sets, LLM-as-judge, groundedness and answer relevance metrics

**Backend:** FastAPI, PostgreSQL, Redis, Celery, REST API design

**Cloud:** Azure (App Service, Blob Storage, Key Vault), Docker, GitHub Actions

## Work Experience

**Machine Learning Engineer - GenAI - Charminar Digital Solutions Pvt Ltd**
*Hyderabad, Telangana | Jan 2024 - Present*

- Designed and shipped a RAG assistant over 12,000 policy and process documents using Qdrant, hybrid BM25 and vector retrieval and a cross-encoder re-ranker; answer acceptance by agents rose from 58% to 81%.
- Built an evaluation pipeline with 250 golden questions and RAGAS metrics that runs on every prompt or retriever change in CI.
- Reduced hallucinated answers by adding retrieval confidence thresholds, citation enforcement and a safe fallback response.
- Cut average LLM cost per query by 35% through context compression, caching of frequent queries in Redis and routing simple queries to a smaller model.
- Exposed the assistant through FastAPI with role-based access controls and audit logging.

**Software Engineer - Charminar Digital Solutions Pvt Ltd**
*Hyderabad, Telangana | Aug 2021 - Dec 2023*

- Built Python microservices and REST APIs for a customer support platform handling about 50,000 tickets per month.
- Implemented a ticket similarity feature using sentence embeddings and pgvector, which became the starting point for the GenAI team.
- Improved PostgreSQL query performance on ticket search, reducing p95 latency from 900 ms to 220 ms.

## Projects

**Contract Clause Q&A**

- Problem: Legal ops needed fast answers from long vendor contracts.
- Approach: Used structure-aware chunking by clause, metadata filters for vendor and date, hybrid retrieval and answer citations.
- Technologies: Python, LlamaIndex, Qdrant, Azure OpenAI
- Results: Reviewers reported saving about 30 minutes per contract query batch.

**RAG Evaluation Dashboard**

- Problem: Teams lacked visibility into retrieval and answer quality over time.
- Approach: Logged queries, retrieved chunks and judge scores, then built a dashboard to surface regressions.
- Technologies: Python, PostgreSQL, Streamlit, RAGAS
- Results: Detected two retrieval regressions before release.

## Education

**B.Tech, Computer Science and Engineering**, Godavari Institute of Technology, Hyderabad (2017 - 2021) - CGPA 7.9 / 10

## Certifications

- Microsoft Azure AI Engineer Associate (AI-102)
- DeepLearning.AI: Building Systems with the ChatGPT API

## Achievements

- Lowered LLM serving cost by 35% while improving acceptance.
- Internal workshop on RAG evaluation attended by 60 engineers.